

- Yoshua Bengio, Geoffrey Hinton, Andrew Yao, Dawn Song, Pieter Abbeel, Yuval Noah Harari, Ya-Qin Zhang, Lan Xue, Shai Shalev-Shwartz, Gillian Hadfield, Jeff Clune, Tegan Maharaj, Frank Hutter, Atılım Güneş Baydin, Sheila McIlraith, Qiqi Gao, Ashwin Acharya, David Krueger, Anca Dragan, Philip Torr, Stuart Russell, Daniel Kahneman, Jan Brauner, and Sören Mindermann. [Managing ai risks in an era of rapid progress](#). *CoRR*, November 2023. 1
- Leonard Bereska. [Mechanistic interpretability for adversarial robustness — a proposal](#). *Leonard Bereska’s Blog*, August 2024. 28
- Leonard Bereska and Efstratios Gavves. [Taming simulators: Challenges, pathways and vision for the alignment of large language models](#). *AAAI-SS*, October 2023. 24, 30, 34
- Michel Besserve, Arash Mehrjou, Rémy Sun, and Bernhard Schölkopf. [Counterfactuals uncover the modular structure of deep generative models](#). *CoRR*, December 2019. 28
- Steven Bills, Nick Cammarata, Dan Mossing, Henk Tillman, Leo Gao, Gabriel Goh, Ilya Sutskever, Jan Leike, Jeff Wu, and William Saunders. [Language models can explain neurons in language models](#). *OpenAI Blog*, 2023. 5, 23, 25
- Blair Bilodeau, Natasha Jaques, Pang Wei Koh, and Been Kim. [Impossibility theorems for feature attribution](#). *Proc. Natl. Acad. Sci. U.S.A.*, January 2024. 2
- Christopher M. Bishop. [Pattern recognition and machine learning](#). Springer-Verlag New York Inc., 2006. 4
- Sid Black, Lee Sharkey, Leo Grinsztajn, Eric Winsor, Dan Braun, Jacob Merizian, Kip Parker, Carlos Ramón Guevara, Beren Millidge, Gabriel Alfour, and Connor Leahy. [Interpreting neural networks through the polytope lens](#). *CoRR*, November 2022. 9, 28
- Joseph Bloom and Jay Bailey. [Features and adversaries in memorydt](#). *LessWrong*, October 2023. 29, 30
- Joseph Bloom and Paul Colognese. [Decision transformer interpretability](#). *AI Alignment Forum*, 2023. 30
- Tolga Bolukbasi, Adam Pearce, Ann Yuan, Andy Coenen, Emily Reif, Fernanda Vi’egas, and M. Wattenberg. [An interpretability illusion for bert](#). *CoRR*, April 2021. 29
- Dan Braun, Jordan Taylor, Nicholas Goldowsky-Dill, and Lee Sharkey. [Identifying functionally important features with end-to-end sparse dictionary learning](#). *ICML MI Workshop*, May 2024. 19
- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nicholas L. Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E. Burke, Tristan Hume, Shan Carter, Tom Henighan, and Chris Olah. [Towards monosemanticity: Decomposing language models with dictionary learning](#). *Transformer Circuits Thread*, October 2023. 5, 6, 7, 8, 9, 11, 13, 15, 20, 30, 32
- Jannik Brinkmann, Abhay Sheshadri, Victor Levoso, Paul Swoboda, and Christian Bartelt. [A mechanistic analysis of a transformer trained on a symbolic multi-step reasoning task](#). *CoRR*, February 2024. 22
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, Harsha Nori, Hamid Palangi, Marco Tulio Ribeiro, and Yi Zhang. [Sparks of artificial general intelligence: Early experiments with gpt-4](#). *CoRR*, April 2023. 1
- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. [Discovering latent knowledge in language models without supervision](#). *ICLR*, 2023. 3, 13, 14, 24, 29
- Lucius Bushnaq, Stefan Heimersheim, Nicholas Goldowsky-Dill, Dan Braun, Jake Mendel, Kaarel Hanni, Avery Griffin, Jorn Stohler, Magdalena Wache, and Marius Hobbhahn. [The local interaction basis: Identifying computationally-relevant and sparsely interacting features in neural networks](#). *CoRR*, May 2024. 19

- Ethan Caballero, Kshitij Gupta, Irina Rish, and David Krueger. [Broken neural scaling laws](#). *ICLR*, October 2022. 21
- Nick Cammarata, Gabriel Goh, Shan Carter, Ludwig Schubert, Michael Petrov, and Chris Olah. [Curve detectors](#). *Distill*, June 2020. 10, 15, 21, 30, 32
- Nick Cammarata, Gabriel Goh, Shan Carter, Chelsea Voss, Ludwig Schubert, and Chris Olah. [Curve circuits](#). *Distill*, 2021. 10, 21, 30, 32
- Steven Cao, Victor Sanh, and Alexander M. Rush. [Low-complexity probing via finding subnetworks](#). *NAACL-HLT*, April 2021. 14
- Shan Carter, Zan Armstrong, Ludwig Schubert, Ian Johnson, and Chris Olah. [Activation atlas](#). *Distill*, March 2019. 28
- Giuseppe Casalicchio, Christoph Molnar, and Bernd Bischl. [Visualizing the feature importance for black box models](#). *ECML PKDD*, 2018. 2
- Stephen Casper. [The engineer’s interpretability sequence](#). *AI Alignment Forum*, February 2023. 23, 25, 27, 28, 29
- Stephen Casper, Max Nadeau, Dylan Hadfield-Menell, and Gabriel Kreiman. [Robust feature-level adversaries are interpretability tools](#). *NeurIPS*, October 2021. 28
- Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémy Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, Tony Wang, Samuel Marks, Charbel-Raphaël Segerie, Micah Carroll, Andi Peng, Phillip Christoffersen, Mehul Damani, Stewart Slocum, Usman Anwar, Anand Siththaranjan, Max Nadeau, Eric J. Michaud, Jacob Pfau, Dmitrii Krashennnikov, Xin Chen, Lauro Langosco, Peter Hase, Erdem Biyık, Anca Dragan, David Krueger, Dorsa Sadigh, and Dylan Hadfield-Menell. [Open problems and fundamental limitations of reinforcement learning from human feedback](#). *CoRR*, 2023a. 24, 30
- Stephen Casper, Kaivalya Hariharan, and Dylan Hadfield-Menell. [Diagnostics for deep neural networks with automated copy/paste attacks](#). *NeurIPS 2022 ML Safety Workshop (Best paper award)*, May 2023b. 25
- Stephen Casper, Yuxiao Li, Jiawei Li, Tong Bu, Kevin Zhang, Kaivalya Hariharan, and Dylan Hadfield-Menell. [Red teaming deep neural networks with feature synthesis tools](#). *NeurIPS*, 2023c. 25, 29
- Stephen Casper, Carson Ezell, Charlotte Siegmann, Noam Kolt, Taylor Lynn Curtis, Benjamin Bucknall, Andreas Haupt, Kevin Wei, Jérémy Scheurer, Marius Hobbhahn, Lee Sharkey, Satyapriya Krishna, Marvin Von Hagen, Silas Alberti, Alan Chan, Qinyi Sun, Michael Gerovitch, David Bau, Max Tegmark, David Krueger, and Dylan Hadfield-Menell. [Black-box access is insufficient for rigorous ai audits](#). *ACM Conference on Fairness, Accountability, and Transparency*, January 2024. 1, 28
- Lawrence Chan. [What i would do if i wasn’t at arc evals](#). *AI Alignment Forum*, May 2023. 24, 28, 29
- Lawrence Chan, Adrià Garriga-alonso, Nicholas Goldowsky-Dill, ryan_greenblatt, jenny, Ansh Radhakrishnan, Buck, and Nate Thomas. [Causal scrubbing: a method for rigorously testing interpretability hypotheses \[redwood research\]](#). *AI Alignment Forum*, December 2022. 13, 19
- Lawrence Chan, Leon Lang, and Erik Jenner. [Natural abstractions: Key claims, theorems, and critiques](#). *AI Alignment Forum*, March 2023. 4, 10, 33
- David Chanin, Anthony Hunter, and Oana-Maria Camburu. [Identifying linear relational concepts in large language models](#). *CoRR*, 2023. 9
- Charbel-Raphaël. [Against almost every theory of impact of interpretability](#). *AI Alignment Forum*, August 2023. 25

- Yiting Chen, Zhanpeng Zhou, and Junchi Yan. [Going beyond neural network feature similarity: The network feature complexity and its interpretation using category theory](#). *CoRR*, November 2023a. 11
- Zhongtian Chen, Edmund Lau, Jake Mendel, Susan Wei, and Daniel Murfet. [Dynamical versus bayesian phase transitions in a toy model of superposition](#). *CoRR*, October 2023b. 21, 25
- Paul Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. [Deep reinforcement learning from human preferences](#). *NeurIPS*, December 2017. 30
- Paul Christiano, Ajeya Cotra, and Mark Xu. [Eliciting latent knowledge](#), January 2021. 24, 32
- Bilal Chughtai, Lawrence Chan, and Neel Nanda. [A toy model of universality: Reverse engineering how networks learn group operations](#). *ICML*, 2023. 11, 22, 25
- Bilal Chughtai, Alan Cooney, and Neel Nanda. [Summing up the facts: Additive mechanisms behind factual recall in llms](#). *NeurIPS Workshop Attributing Model Behaviour at Scale*, 2024. 10
- Paul Colognese. [Internal target information for ai oversight](#). *LessWrong*, 2023. 30
- Paul Colognese and Jozdien. [High-level interpretability: detecting an ai’s objectives](#). *AI Alignment Forum*, 2023. 30
- Arthur Conmy, Augustine N. Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. [Towards automated circuit discovery for mechanistic interpretability](#). *NeurIPS*, 2023. 18, 19, 22, 25, 30
- Ian C. Covert, Scott Lundberg, and Su-In Lee. [Explaining by removing: a unified framework for model explanation](#). *J. Mach. Learn. Res.*, January 2021. 2
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. [Sparse autoencoders find highly interpretable features in language models](#). *ICLR*, January 2024. 8, 9, 13, 15, 16, 19, 29
- Damai Dai, Li Dong, Yaru Hao, Zhifang Sui, Baobao Chang, and Furu Wei. [Knowledge neurons in pretrained transformers](#). *ACL*, 2022. 5, 29
- David "davidad" Dalrymple, Joar Skalse, Yoshua Bengio, Stuart Russell, Max Tegmark, Sanjit Seshia, Steve Omohundro, Christian Szegedy, Ben Goldhaber, Nora Ammann, Alessandro Abate, Joe Halpern, Clark Barrett, Ding Zhao, Tan Zhi-Xuan, Jeannette Wing, and Joshua Tenenbaum. [Towards guaranteed safe ai: A framework for ensuring robust and reliable ai systems](#). *CoRR*, May 2024. 29, 35
- Fahim Dalvi, Nadir Durrani, Hassan Sajjad, Yonatan Belinkov, Anthony Bau, and James Glass. [What is one grain of sand in the desert? analyzing individual neurons in deep nlp models](#). *Proceedings of the AAAI Conference on Artificial Intelligence*, July 2019. 14
- Guy Dar, Mor Geva, Ankit Gupta, and Jonathan Berant. [Analyzing transformers in embedding space](#). *ACL*, December 2022. 15
- Xander Davies, Max Nadeau, Nikhil Prakash, Tamar Rott Shaham, and David Bau. [Discovering variable binding circuitry with desiderata](#). *CoRR*, July 2023. 10, 22
- Mingyang Deng, Lucas Tao, and Joe Benton. [Measuring feature sparsity in language models](#). *CoRR*, 2023. 9, 15
- Carson Denison, Monte MacDiarmid, Fazl Barez, David Duvenaud, Shauna Kravec, Samuel Marks, Nicholas Schiefer, Ryan Soklaski, Alex Tamkin, Jared Kaplan, Buck Shlegeris, Samuel R. Bowman, Ethan Perez, and Evan Hubinger. [Sycophancy to subterfuge: Investigating reward-tampering in large language models](#). *CoRR*, 2024. 28
- Alexander Yom Din, Taelin Karidi, Leshem Choshen, and Mor Geva. [Jump to conclusions: Short-cutting transformers with linear transformations](#). *CoRR*, March 2023. 15