

- Finale Doshi-Velez and Been Kim. [Towards a rigorous science of interpretable machine learning](#). *CoRR*, March 2017. 21, 29, 30
- Keke Du, Shan Chang, Huixiang Wen, and Hao Zhang. [Fighting adversarial images with interpretable gradients](#). *ACM TURC*, October 2021. 28
- Jacob Dunefsky, Philippe Chlenski, and Neel Nanda. [Transcoders find interpretable llm feature circuits](#). *ICML MI Workshop*, June 2024. 16
- Nadir Durrani, Hassan Sajjad, Fahim Dalvi, and Yonatan Belinkov. [Analyzing individual neurons in pre-trained language models](#). *EMNLP*, October 2020. 5
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. [Amnesic probing: Behavioral explanation with amnesic counterfactuals](#). *TACL*, February 2021. 14
- Nelson Elhage, Tristan Hume, Olsson Catherine, Nanda Neel, Tom Henighan, Scott Johnston, Sheer ElShowk, Nicholas Joseph, Nova DasSarma, Ben Mann, Danny Hernandez, Amanda Askell, Kamal Ndousse, Dawn Drain, Anna Chen, Yuntao Bai, Deep Ganguli, Liane Lovitt, Zac Hatfield-Dodds, Jackson Kernion, Tom Conerly, Shauna Kravec, Stanislav Fort, Saurav Kadavath, Josh Jacobson, Eli Tran-Johnson, Jared Kaplan, Jack Clark, Tom Brown, Sam McCandlish, Dario Amodei, and Christopher Olah. [Softmax linear units](#). *Transformer Circuits Thread*, 2022a. 5, 6, 13, 20
- Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, et al. [Toy models of superposition](#). *Transformer Circuits Thread*, 2022b. 5, 6, 7, 9, 15, 25, 26, 27, 28, 29, 30, 34, 35
- Joshua Engels, Isaac Liao, Eric J. Michaud, Wes Gurnee, and Max Tegmark. [Not all language model features are linear](#). *CoRR*, May 2024. 4, 8, 33
- Logan Engstrom, Andrew Ilyas, Shibani Santurkar, Dimitris Tsipras, Brandon Tran, and Aleksander Madry. [Adversarial robustness as a prior for learned representations](#). *CoRR*, September 2019. 28
- Ingrid Lossius Falkum and Agustin Vicente. [Polysemy: Current perspectives and approaches](#). *Lingua*, April 2015. 34
- Sebastian Farquhar, Vikrant Varma, Zachary Kenton, Johannes Gasteiger, Vladimir Mikulik, and Rohin Shah. [Challenges with unsupervised llm knowledge discovery](#). *CoRR*, 2023. 14
- Amir Feder, Nadav Oved, Uri Shalit, and Roi Reichart. [Causalm: Causal model explanation through counterfactual language models](#). *Computational Linguistics*, May 2021. 18
- Jiahai Feng and Jacob Steinhardt. [How do language models bind entities in context?](#) *CoRR*, October 2023. 10, 22, 26, 30
- Javier Ferrando and Elena Voita. [Information flow routes: Automatically interpreting language models at scale](#). *CoRR*, February 2024. 18, 22
- Javier Ferrando, Gabriele Sarti, Arianna Bisazza, and Marta R. Costa-jussà. [A primer on the inner workings of transformer-based language models](#). *CoRR*, May 2024. 2
- Alex Foote, Neel Nanda, Esben Kran, Ioannis Konstas, Shay Cohen, and Fazl Barez. [Neuron to graph: Interpreting language model neurons at scale](#). *CoRR*, May 2023. 23, 30
- Stanislav Fort and Balaji Lakshminarayanan. [Ensemble everything everywhere: Multi-scale aggregation for adversarial robustness](#). *CoRR*, August 2024. 28
- Jonathan Frankle and Michael Carbin. [The lottery ticket hypothesis: Finding sparse, trainable neural networks](#). *ICLR*, March 2019. 20
- Dan Friedman, Andrew Lampinen, Lucas Dixon, Danqi Chen, and Asma Ghandeharioun. [Interpretability illusions in the generalization of simplified models](#). *CoRR*, 2023a. 29, 30

- Dan Friedman, Alexander Wettig, and Danqi Chen. [Learning transformer programs](#). *NeurIPS*, June 2023b. 20
- Daniel Y. Fu, Tri Dao, Khaled K. Saab, Armin W. Thomas, Atri Rudra, and Christopher Ré. [Hungry hungry hippos: Towards language modeling with state space models](#). *ICLR*, 2023a. 25
- Deqing Fu, Tian-Qi Chen, Robin Jia, and Vatsal Sharan. [Transformers learn higher-order optimization methods for in-context learning: A study with linear models](#). *CoRR*, October 2023b. 26
- Zach Furman and Edmund Lau. [Estimating the local learning coefficient at scale](#). *CoRR*, February 2024. 21
- Jorge García-Carrasco, Alejandro Maté, and Juan Trujillo. [Detecting and understanding vulnerabilities in language models via mechanistic interpretability](#). *IJCAI*, August 2024. 28, 29
- Albert Garde, Esben Kran, and Fazl Barez. [Deepdecipher: Accessing and investigating neuron activation in large language models](#). *NeurIPS Workshop XAIA*, October 2023. 13
- Peter Gardenfors. *Conceptual spaces: The geometry of thought*. MIT press, 2004. 12
- Charles J. Garfinkle and Christopher J. Hillar. [On the uniqueness and stability of dictionaries for sparse representation of noisy signals](#). *IEEE Transactions on Signal Processing*, December 2019. 15
- Xuyang Ge, Fukang Zhu, Wentao Shu, Junxuan Wang, Zhengfu He, and Xipeng Qiu. [Automatically identifying local and global circuits with linear computation graphs](#). *CoRR*, May 2024. 19
- Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. [Causal abstractions of neural networks](#). *NeurIPS*, 2021a. 18
- Atticus Geiger, Zhengxuan Wu, Hanson Lu, Josh Rozner, Elisa Kreiss, Thomas Icard, Noah D. Goodman, and Christopher Potts. [Inducing causal structure for interpretable neural networks](#). *ICML*, January 2021b. 17
- Atticus Geiger, Chris Potts, and Thomas Icard. [Causal abstraction for faithful model interpretation](#). *CoRR*, January 2023a. 13, 16, 18, 19
- Atticus Geiger, Zhengxuan Wu, Christopher Potts, Thomas Icard, and Noah D. Goodman. [Finding alignments between interpretable causal variables and distributed neural representations](#). *CoRR*, 2023b. 13, 18
- Georgios Georgiadis. [Accelerating convolutional neural networks via activation map compression](#). *CoRR*, March 2019. 20
- Mor Geva, Avi Caciularu, Kevin Ro Wang, and Yoav Goldberg. [Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space](#). *EMNLP*, October 2022. 15
- Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. [Dissecting recall of factual associations in auto-regressive language models](#). *EMNLP*, October 2023. 17
- Asma Ghandeharioun, Avi Caciularu, Adam Pearce, Lucas Dixon, and Mor Geva. [Patchscopes: A unifying framework for inspecting hidden representations of language models](#). *CoRR*, January 2024. 18
- Amirata Ghorbani and James Zou. [Neuron shapley: Discovering the responsible neurons](#). *NeurIPS*, November 2020. 5, 29
- Gabriel Goh, Nick Cammarata †, Chelsea Voss †, Shan Carter, Michael Petrov, Ludwig Schubert, Alec Radford, and Chris Olah. [Multimodal neurons in artificial neural networks](#). *Distill*, March 2021. 5, 30
- Nicholas Goldowsky-Dill, Chris MacLeod, Lucas Sato, and Aryaman Arora. [Localizing model behavior with path patching](#). *CoRR*, 2023. 13, 17, 18, 29, 30
- Liv Gorton. [The missing curve detectors of inceptionv1: Applying sparse autoencoders to inceptionv1 early vision](#). *ICML MI Workshop*, June 2024. 15

- Rhys Gould, Euan Ong, George Ogden, and Arthur Conmy. [Successor heads: Recurring, interpretable attention heads in the wild](#). *CoRR*, 2023. 29
- Anirudh Goyal, Alex Lamb, Jordan Hoffmann, Shagun Sodhani, Sergey Levine, Yoshua Bengio, and Bernhard Schölkopf. [Recurrent independent mechanisms](#). *CoRR*, November 2020. 28
- Jason Gross, Rajashree Agrawal, Thomas Kwa, Euan Ong, Chun Hei Yip, Alex Gibson, Soufiane Noubir, and Lawrence Chan. [Compact proofs of model performance via mechanistic interpretability](#). *ICML MI Workshop*, June 2024. 29
- Roger Grosse, Juhan Bae, Cem Anil, Nelson Elhage, Alex Tamkin, Amirhossein Tajdini, Benoit Steiner, Dustin Li, Esin Durmus, Ethan Perez, Evan Hubinger, Kamilė Lukošiuūtė, Karina Nguyen, Nicholas Joseph, Sam McCandlish, Jared Kaplan, and Samuel R. Bowman. [Studying large language model generalization with influence functions](#). *CoRR*, August 2023. 13
- Phillip Huang Guo, Aaquib Syed, Abhay Sheshadri, Aidan Ewart, and Gintare Karolina Dziugaite. [Robust unlearning via mechanistic localizations](#). *ICML MI Workshop*, June 2024. 25
- Rohan Gupta, Iván Arcuschin, Thomas Kwa, and Adrià Garriga-Alonso. [Interpbench: Semi-synthetic transformers for evaluating mechanistic interpretability techniques](#). *CoRR*, July 2024. 29
- Wes Gurnee and Max Tegmark. [Language models represent space and time](#). *ICLR*, 2024. 12
- Wes Gurnee, Neel Nanda, Matthew Pauly, Katherine Harvey, Dmitrii Troitskii, and Dimitris Bertsimas. [Finding neurons in a haystack: Case studies with sparse probing](#). *TMLR*, 2023. 6, 13, 14
- Wes Gurnee, Theo Horsley, Zifan Carl Guo, Tara Rezaei Kheirkhah, Qinyi Sun, Will Hathaway, Neel Nanda, and Dimitris Bertsimas. [Universal neurons in gpt2 language models](#). *CoRR*, January 2024. 11
- David R. Ha and J. Schmidhuber. [Recurrent world models facilitate policy evolution](#). *NeurIPS*, September 2018. 11
- Guy Hacoheh, Leshem Choshen, and Daphna Weinshall. [Let’s agree to agree: Neural networks share classification order on real datasets](#). *ICML*, 2020. 11
- Michael Hanna, Ollie Liu, and Alexandre Variengien. [How does gpt-2 compute greater-than?: Interpreting mathematical abilities in a pre-trained language model](#). *NeurIPS*, 2023. 10, 17, 22, 30
- Michael Hanna, Sandro Pezzelle, and Yonatan Belinkov. [Have faith in faithfulness: Going beyond circuit overlap when finding model mechanisms](#). *ICML MI Workshop*, June 2024. 29
- Kaarel Hänni, Jake Mendel, Dmitry Vaintrob, and Lawrence Chan. [Mathematical models of computation in superposition](#). *ICML MI Workshop*, August 2024. 29
- Peter Hase, Mohit Bansal, Been Kim, and Asma Ghandeharioun. [Does localization inform editing? surprising differences in causality-based localization vs. knowledge editing in language models](#). *NeurIPS Spotlight*, January 2023. 29
- Zhengfu He, Xuyang Ge, Qiong Tang, Tianxiang Sun, Qinyuan Cheng, and Xipeng Qiu. [Dictionary learning improves patch-free circuit discovery in mechanistic interpretability: A case study on othello-gpt](#). *CoRR*, 2024. 15
- Stefan Heimersheim and Jett. [A circuit for python docstrings in a 4-layer attention-only transformer](#). *AI Alignment Forum*, February 2023. 22
- Roe Hendel, Mor Geva, and Amir Globerson. [In-context learning creates task vectors](#). *EMNLP*, October 2023. 9, 17, 22
- Dan Hendrycks. [Introduction to AI Safety, Ethics, and Society](#). Self-published, 2023a. 26
- Dan Hendrycks. [Natural selection favors ais over humans](#). *CoRR*, July 2023b. 24