

- Dan Hendrycks and Mantas Mazeika. [X-risk analysis for ai research](#). *CoRR*, June 2022. 1, 24, 25
- Dan Hendrycks, Nicholas Carlini, John Schulman, and Jacob Steinhardt. [Unsolved problems in ml safety](#). *CoRR*, June 2022. 26
- Dan Hendrycks, Mantas Mazeika, and Thomas Woodside. [An overview of catastrophic ai risks](#). *CoRR*, October 2023. 1, 24
- Tom Henighan, Shan Carter, Tristan Hume, Nelson Elhage, Robert Lasenby, Stanislaw Fort, Nicholas Schiefer, and Christopher Olah. [Superposition, memorization, and double descent](#). *Transformer Circuits Thread*, 2023. 6, 27
- Danny Hernandez, Tom Brown, Tom Conerly, Nova DasSarma, Dawn Drain, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Tom Henighan, Tristan Hume, Scott Johnston, Ben Mann, Chris Olah, Catherine Olsson, Dario Amodei, Nicholas Joseph, Jared Kaplan, and Sam McCandlish. [Scaling laws and interpretability of learning from repeated data](#). *CoRR*, 2022. 21
- Evan Hernandez, Arnab Sen Sharma, Tal Haklay, Kevin Meng, Martin Wattenberg, Jacob Andreas, Yonatan Belinkov, and David Bau. [Linearity of relation decoding in transformer language models](#). *CoRR*, August 2023. 9
- John Hewitt and Christopher D. Manning. [A structural probe for finding syntax in word representations](#). *NAACL HLT*, June 2019. 14
- Irina Higgins, David Amos, David Pfau, Sebastien Racaniere, Loic Matthey, Danilo Rezende, and Alexander Lerchner. [Towards a definition of disentangled representations](#). *CoRR*, December 2018. 27
- Jacob Hilton, Nick Cammarata, Shan Carter, Gabriel Goh, and Chris Olah. [Understanding rl vision](#). *Distill*, 2020. 30
- Geoffrey E Hinton. [Distributed representations](#). *Carnegie Mellon University*, 1984. 27
- Marius Hobbhahn. [Marius’ alignment agenda](#), 2022. 30
- Marius Hobbhahn and Lawrence Chan. [Should we publish mechanistic interpretability research?](#) *AI Alignment Forum*, April 2023. 25
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katie Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. [Training compute-optimal large language models](#). *CoRR*, March 2022. 25
- Jesse Hoogland, Liam Carroll, and Daniel Murfet. [Stagewise development in neural networks](#). *AI Alignment Forum*, March 2024. 21
- Jing Huang, Atticus Geiger, Karel D’Oosterlinck, Zhengxuan Wu, and Christopher Potts. [Rigorously assessing natural language explanations of neurons](#). *CoRR*, September 2023. 5
- Jing Huang, Zhengxuan Wu, Christopher Potts, Mor Geva, and Atticus Geiger. [Ravel: Evaluating interpretability methods on disentangling language model representations](#). *CoRR*, 2024. 19
- Evan Hubinger. [Chris olah’s views on agi safety](#). *AI Alignment Forum*, November 2019a. 4, 24, 26
- Evan Hubinger. [Gradient hacking](#). *AI Alignment Forum*, October 2019b. 26
- Evan Hubinger. [An overview of 11 proposals for building safe advanced ai](#). *CoRR*, December 2020. 24
- Evan Hubinger. [A transparency and interpretability tech tree](#). *AI Alignment Forum*, June 2022. 9, 24, 26, 30

- Evan Hubinger, Chris van Merwijk, Vladimir Mikulik, Joar Skalse, and Scott Garrabrant. [Risks from learned optimization in advanced machine learning systems](#). *CoRR*, May 2019. 24, 31
- Evan Hubinger, Adam Jermy, Johannes Treutlein, Rubi Hudson, and Kate Woolverton. [Conditioning predictive models: Risks and strategies](#). *CoRR*, February 2023. 24, 31
- Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tamera Lanham, Daniel M. Ziegler, Tim Maxwell, Newton Cheng, Adam Jermy, Amanda Askill, Ansh Radhakrishnan, Cem Anil, David Duvenaud, Deep Ganguli, Fazl Barez, Jack Clark, Kamal Ndousse, Kshitij Sachan, Michael Sellitto, Mrinank Sharma, Nova DasSarma, Roger Grosse, Shauna Kravec, Yuntao Bai, Zachary Witten, Marina Favaro, Jan Brauner, Holden Karnofsky, Paul Christiano, Samuel R. Bowman, Logan Graham, Jared Kaplan, Sören Mindermann, Ryan Greenblatt, Buck Shlegeris, Nicholas Schiefer, and Ethan Perez. [Sleeper agents: Training deceptive llms that persist through safety training](#). *CoRR*, 2024. 28, 29
- Andrew Ilyas, Shibani Santurkar, Dimitris Tsipras, Logan Engstrom, Brandon Tran, and Aleksander Madry. [Adversarial examples are not bugs, they are features](#). *NeurIPS*, August 2019. 4, 28, 29
- M. Ivanitskiy, Alexander F. Spies, Tilman Rauker, Guillaume Corlourer, Chris Mathwin, Lucia Quirke, Can Rager, Rusheb Shah, Dan Valentine, Cecilia Diniz Behn, Katsumi Inoue, and Samy Wu Fung. [Structured world representations in maze-solving transformers](#). *CoRR*, December 2023. 12, 30, 31
- Alon Jacovi and Yoav Goldberg. [Towards faithfully interpretable nlp systems: How should we define and evaluate faithfulness?](#) *CoRR*, April 2020. 29
- Samyak Jain, Robert Kirk, Ekdeep Singh Lubana, Robert P. Dick, Hidenori Tanaka, Edward Grefenstette, Tim Rocktäschel, and David Scott Krueger. [Mechanistically analyzing the effects of fine-tuning on procedurally defined tasks](#). *CoRR*, November 2023. 30
- Samyak Jain, Ekdeep Singh Lubana, Kemal Oksuz, Tom Joy, Philip Torr, Amartya Sanyal, and Puneet K. Dokania. [What makes and breaks safety fine-tuning? a mechanistic study](#). *ICML MI Workshop*, June 2024. 30
- janus. [Simulators](#). *LessWrong*, September 2022. 11, 12, 24, 34
- Erik Jenner, Adrià Garriga-alonso, and Egor Zverev. [A comparison of causal scrubbing, causal abstractions, and related methods](#). *AI Alignment Forum*, June 2023. 13, 18, 19
- Erik Jenner, Shreyas Kapur, Vasil Georgiev, Cameron Allen, Scott Emmons, and Stuart Russell. [Evidence of learned look-ahead in a chess-playing neural network](#). *CoRR*, June 2024. 24
- Adam Jermy, Chris Olah, and T Henighan. [Circuits updates - may 2023: Attention head superposition](#). *Transformer Circuits Thread*, 2023. 29
- Adam S. Jermy, Nicholas Schiefer, and Evan Hubinger. [Engineering monosemanticity in toy models](#). *CoRR*, November 2022. 6, 20, 25
- Jiaming Ji, Tianyi Qiu, Boyuan Chen, Borong Zhang, Hantao Lou, Kaile Wang, Yawen Duan, Zhonghao He, Jiayi Zhou, Zhaowei Zhang, Fanzhi Zeng, Kwan Yee Ng, Juntao Dai, Xuehai Pan, Aidan O’Gara, Yingshan Lei, Hua Xu, Brian Tse, Jie Fu, Stephen McAleer, Yaodong Yang, Yizhou Wang, Song-Chun Zhu, Yike Guo, and Wen Gao. [Ai alignment: A comprehensive survey](#). *CoRR*, January 2024. 1
- Jozdian. [Conditioning generative models for alignment](#). *AI Alignment Forum*, July 2022. 24
- Jaap Jumelet. [Evaluating and interpreting language models](#). *NLP Lecture*, November 2023. 2
- Amlan Jyoti, Karthik Balaji Ganesh, Manoj Gayala, Nandita Lakshmi Tunuguntla, Sandesh Kamath, and Vineeth N. Balasubramanian. [On the robustness of explanations of deep neural network models: A survey](#). *CoRR*, November 2022. 28

- Adam Karvonen. [Emergent world models and latent variable estimation in chess-playing language models.](#) *COLM*, July 2024. 11
- Adam Karvonen, Benjamin Wright, Can Rager, Rico Angell, Jannik Brinkmann, Logan Smith, C. M. Verdun, David Bau, and Samuel Marks. [Measuring progress in dictionary learning for language model interpretability with board game models.](#) *ICML MI Workshop (Oral)*, July 2024. 15
- Theodoros Kasioumis, Joe Townsend, and Hiroya Inakoshi. [Elite backprop: Training sparse interpretable neurons.](#) *International Workshop on Neuro-Symbolic Learning and Reasoning*, 2021. 20
- Nan Rosemary Ke, Aniket Didolkar, Sarthak Mittal, Anirudh Goyal, Guillaume Lajoie, Stefan Bauer, Danilo Rezende, Yoshua Bengio, Michael Mozer, and Christopher Pal. [Systematic evaluation of causal discovery in visual model based reinforcement learning.](#) *CoRR*, July 2021. 28
- Jan Kirchner. [Neuroscience and natural abstractions.](#) *LessWrong*, March 2023. 11
- Connor Kissane, Robert Krzyzanowski, Joseph Isaac Bloom, Arthur Conmy, and Neel Nanda. [Interpreting attention layer outputs with sparse autoencoders.](#) *ICML MI Workshop*, June 2024. 15
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. [Similarity of neural network representations revisited.](#) *ICML*, July 2019. 3, 10
- János Kramár, Tom Lieberum, Rohin Shah, and Neel Nanda. [Atp*: An efficient and scalable method for localizing llm behaviour to components.](#) *CoRR*, March 2024. 22
- Nicholas Kross. [Why and when interpretability work is dangerous.](#) *LessWrong*, May 2023. 25
- Jan Kulveit. [Risks from ai misalignment at different scales](#), July 2024. 24
- Jan Kulveit, Clem von Stengel, and Roman Leventov. [Predictive minds: Llms as atypical active inference agents.](#) *CoRR*, November 2023. 11
- Max Lamparth and Anka Reuel. [Analyzing and editing inner mechanisms of backdoored language models.](#) *CoRR*, 2023. 29
- Michael Lan and Fazl Barez. [Locating cross-task sequence continuation circuits in transformers.](#) *CoRR*, November 2023. 10
- Leon Lang, Davis Foote, Stuart Russell, Anca Dragan, Erik Jenner, and Scott Emmons. [When your ais deceive you: Challenges with partial observability of human evaluators in reward learning.](#) *CoRR*, March 2024. 28, 32
- Georg Lange, Alex Makelov, and Neel Nanda. [An interpretability illusion for activation patching of arbitrary subspaces.](#) *AI Alignment Forum*, August 2023. 18
- Anna Langedijk, Hosein Mohebbi, Gabriele Sarti, Willem Zuidema, and Jaap Jumelet. [Decoderlens: Layer-wise interpretation of encoder-decoder transformers.](#) *CoRR*, 2023. 15
- Edmund Lau, Daniel Mufet, and Susan Wei. [Quantifying degeneracy in singular models via the learning coefficient.](#) *CoRR*, August 2023. 21
- Connor Leahy. [Barriers to mechanistic interpretability for agi safety.](#) *AI Alignment Forum*, 2023. 26
- Matthew L. Leavitt and Ari Morcos. [Towards falsifiable interpretability research.](#) *CoRR*, October 2020. 29
- Victor Lecomte, Kushal Thaman, Trevor Chow, Rylan Schaeffer, and Sanmi Koyejo. [Incidental polyseman-
ticity.](#) *CoRR*, 2023. 27, 28, 29
- Andrew Lee, Xiaoyan Bai, Itamar Pres, Martin Wattenberg, Jonathan K. Kummerfeld, and Rada Mihalcea. [A mechanistic understanding of alignment algorithms: A case study on dpo and toxicity.](#) *CoRR*, 2024. 30