

- Belinda Z. Li, Maxwell Nye, and Jacob Andreas. [Implicit representations of meaning in neural language models](#). *ACL-IJCNLP*, August 2021. 12
- Kenneth Li, Aspen K. Hopkins, David Bau, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. [Emergent world representations: Exploring a sequence model trained on a synthetic task](#). *ICLR*, 2023a. 8, 11
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. [Inference-time intervention: Eliciting truthful answers from a language model](#). *NeurIPS Spotlight*, July 2023b. 9
- Yixuan Li, Jason Yosinski, Jeff Clune, Hod Lipson, and John Hopcroft. [Convergent learning: Do different neural networks learn the same representations?](#) *NIPS Workshop on Feature Extraction*, December 2015. 11
- Tom Lieberum, Matthew Rahtz, János Kramár, Neel Nanda, Geoffrey Irving, Rohin Shah, and Vladimir Mikulik. [Does circuit analysis interpretability scale? evidence from multiple choice capabilities in chimilla](#). *CoRR*, July 2023. 10, 17, 22, 29, 30
- David Lindner, János Kramár, Sebastian Farquhar, Matthew Rahtz, Thomas McGrath, and Vladimir Mikulik. [Tracr: Compiled transformers as a laboratory for interpretability](#). *CoRR*, 2023. 29
- Leo Z. Liu, Yizhong Wang, Jungo Kasai, Hannaneh Hajishirzi, and Noah A. Smith. [Probing across time: What does roberta know and when?](#) *EMNLP*, September 2021. 30
- Ziming Liu and Max Tegmark. [A neural scaling law from lottery ticket ensembling](#). *CoRR*, 2023. 21
- Ziming Liu, Ouail Kitouni, Niklas Nolte, Eric J. Michaud, Max Tegmark, and Mike Williams. [Towards understanding grokking: An effective theory of representation learning](#). *NeurIPS*, 2022a. 21, 32
- Ziming Liu, Eric J. Michaud, and Max Tegmark. [Omnigrok: Grokking beyond algorithmic data](#). *ICML*, 2022b. 21
- Ziming Liu, Eric Gan, and Max Tegmark. [Seeing is believing: Brain-inspired modular training for mechanistic interpretability](#). *Entropy*, June 2023a. 13, 20, 23, 29
- Ziming Liu, Mikail Khona, Ila R. Fiete, and Max Tegmark. [Growing brains: Co-emergence of anatomical and functional modularity in recurrent neural networks](#). *CoRR*, 2023b. 20
- Ziming Liu, Ziqian Zhong, and Max Tegmark. [Grokking as compression: A nonlinear complexity perspective](#). *CoRR*, 2023c. 21
- Francesco Locatello, Stefan Bauer, Mario Lucic, Gunnar Rätsch, Sylvain Gelly, Bernhard Schölkopf, and Olivier Bachem. [Challenging common assumptions in the unsupervised learning of disentangled representations](#). *CoRR*, June 2019. 27
- Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. [The ai scientist: Towards fully automated open-ended scientific discovery](#). *CoRR*, August 2024. 26
- Ang Lv, Yuhan Chen, Kaiyi Zhang, Yulong Wang, Lifeng Liu, Ji-Rong Wen, Jian Xie, and Rui Yan. [Interpreting key mechanisms of factual recall in transformer-based language models](#). *CoRR*, 2024. 10
- Sara Magliacane, Thijs van Ommen, Tom Claassen, Stephan Bongers, Philip Versteeg, and Joris M. Mooij. [Domain adaptation by using causal inference to predict invariant conditional distributions](#). *NeurIPS*, October 2018. 28
- Aleksandar Makelov. [Sparse autoencoders match supervised features for model steering on the ioi task](#). *ICML MI Workshop*, June 2024. 15
- Aleksandar Makelov, George Lange, and Neel Nanda. [Towards principled evaluations of sparse autoencoders for interpretability and control](#). *CoRR*, May 2024. 15

- Alex Mallen and Nora Belrose. [Eliciting latent knowledge from quirky language models](#). *CoRR*, December 2023. 32
- Narek Maloyan, Ekansh Verma, Bulat Nutfullin, and Bislan Ashinov. [Trojan detection in large language models: Insights from the trojan detection challenge](#). *CoRR*, April 2024. 28
- Giovanni Luca Marchetti, Christopher Hillar, Danica Kragic, and Sophia Sanborn. [Harmonics of learning: Universal fourier features emerge in invariant networks](#). *CoRR*, December 2023. 11
- Luke Marks, Amir Abdullah, Luna Mendez, Rauno Arike, Philip Torr, and Fazl Barez. [Interpreting reward models in rlhf-tuned language models using sparse autoencoders](#). *CoRR*, October 2023a. 15
- Luke Marks, Amir Abdullah, Clement Neo, Rauno Arike, Philip Torr, and Fazl Barez. [Beyond training objectives: Interpreting reward model divergence in large language models](#). *CoRR*, October 2023b. 31
- Samuel Marks, Can Rager, Eric J. Michaud, Yonatan Belinkov, David Bau, and Aaron Mueller. [Sparse feature circuits: Discovering and editing interpretable causal graphs in language models](#). *CoRR*, March 2024. 19
- Simon C. Marshall and Jan H. Kirchner. [Understanding polysematicity in neural networks through coding theory](#). *CoRR*, January 2024. 6, 27, 28
- Mantas Mazeika, Andy Zou, Akul Arora, Pavel Pleskov, Dawn Song, Dan Hendrycks, Bo Li, and David Forsyth. [How hard is trojan detection in dnns? fooling detectors with evasive trojans](#). *CoRR*, September 2022. 29
- Callum McDougall, Arthur Conmy, Cody Rushing, Thomas McGrath, and Neel Nanda. [Copy suppression: Comprehensively understanding an attention head](#). *CoRR*, October 2023. 10
- Thomas McGrath, Andrei Kapishnikov, Nenad Tomašev, Adam Pearce, Martin Wattenberg, Demis Hassabis, Been Kim, Ulrich Paquet, and Vladimir Kramnik. [Acquisition of chess knowledge in alphazero](#). *PNAS*, November 2022. 13, 14, 26
- Thomas McGrath, Matthew Rahtz, Janos Kramar, Vladimir Mikulik, and Shane Legg. [The hydra effect: Emergent self-repair in language model computations](#). *CoRR*, July 2023. 18, 26, 28, 32
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. [Locating and editing factual associations in gpt](#). *NeurIPS*, 2022a. 13, 17, 24, 29
- Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. [Mass-editing memory in a transformer](#). *ICLR*, 2022b. 29
- William Merrill, Nikolaos Tsilivis, and Aman Shukla. [A tale of two circuits: Grokking as competition of sparse and dense subnetworks](#). *CoRR*, 2023. 21
- Jack Merullo, Carsten Eickhoff, and Ellie Pavlick. [A mechanism for solving relational tasks in transformer language models](#). *CoRR*, May 2023. 30
- Eric J. Michaud, Ziming Liu, Uzay Girit, and Max Tegmark. [The quantization model of neural scaling](#). *CoRR*, March 2023. 9, 21, 30
- Eric J. Michaud, Isaac Liao, Vedang Lad, Ziming Liu, Anish Mudide, Chloe Loughridge, Zifan Carl Guo, Tara Rezaei Kheirkhah, Mateja Vukelić, and Max Tegmark. [Opening the ai black box: program synthesis via mechanistic interpretability](#). *CoRR*, February 2024. 23, 27
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. [Distributed representations of words and phrases and their compositionality](#). *NeurIPS*, October 2013. 9
- Joseph Miller and Clement Neo. [We found an neuron in gpt-2](#). *AI Alignment Forum*, February 2023. 22

- Tim Miller. [Explanation in artificial intelligence: Insights from the social sciences](#). *Artificial Intelligence*, February 2019. 29
- Ulisse Mini, Peli Grietzer, Mrinank Sharma, Austin Meek, Monte MacDiarmid, and Alexander Matt Turner. [Understanding and controlling a maze-solving policy network](#). *CoRR*, October 2023. 30
- Giovanni Monea, Maxime Peyrard, Martin Josifoski, Vishrav Chaudhary, Jason Eisner, Emre Kiciman, Hamid Palangi, Barun Patra, and Robert West. [A glitch in the matrix? locating and detecting language model grounding with fakepedia](#). *CoRR*, 2023. 30
- Basel Mousi, Nadir Durrani, and Fahim Dalvi. [Can llms facilitate interpretation of pre-trained language models?](#) *EMNLP*, 2023. 23
- Jesse Mu and Jacob Andreas. [Compositional explanations of neurons](#). *NeurIPS*, June 2020. 5, 25, 29
- Aaron Mueller, Jannik Brinkmann, Millicent Li, Samuel Marks, Koyena Pal, Nikhil Prakash, Can Rager, Aruna Sankaranarayanan, Arnab Sen Sharma, Jiuding Sun, Eric Todd, David Bau, and Yonatan Belinkov. [The quest for the right mediator: A history, survey, and theoretical grounding of causal interpretability](#). *CoRR*, August 2024. 17
- Jatin Nainani. [Evaluating brain-inspired modular training in automated circuit discovery for mechanistic interpretability](#). *CoRR*, January 2024. 22, 23
- Preetum Nakkiran, Gal Kaplun, Dimitris Kalimeris, Tristan Yang, Benjamin L. Edelman, Fred Zhang, and Boaz Barak. [Sgd on neural networks learns functions of increasing complexity](#). *NeurIPS*, May 2019. 21
- Neel Nanda. [200 cop in mi: Looking for circuits in the wild](#). *Neel Nanda’s Blog*, 2022a. 30
- Neel Nanda. [200 cop in mi: Analysing training dynamics](#). *Neel Nanda’s Blog*, 2022b. 30
- Neel Nanda. [200 cop in mi: Studying learned features in language models](#). *Neel Nanda’s Blog*, 2022c. 30
- Neel Nanda. [A comprehensive mechanistic interpretability explainer & glossary](#). *Neel Nanda’s Blog*, December 2022d. 2
- Neel Nanda. [A longlist of theories of impact for interpretability](#). *AI Alignment Forum*, March 2022e. 23
- Neel Nanda. [200 cop in mi: Exploring polysemanticity and superposition](#). *Neel Nanda’s Blog*, January 2023a. 29
- Neel Nanda. [200 cop in mi: Techniques, tooling and automation](#). *Neel Nanda’s Blog*, 2023b. 30
- Neel Nanda. [Actually, othello-gpt has a linear emergent world representation](#). *Neel Nanda’s Blog*, March 2023c. 8
- Neel Nanda. [Attribution patching: Activation patching at industrial scale](#). *Neel Nanda’s Blog*, February 2023d. 18
- Neel Nanda. [How to think about activation patching](#). *AI Alignment Forum*, April 2023e. 18, 21
- Neel Nanda. [Mechanistic interpretability quickstart guide](#). *Neel Nanda’s Blog*, January 2023f. 2, 30
- Neel Nanda. [An extremely opinionated annotated list of my favourite mechanistic interpretability papers v2](#). *AI Alignment Forum*, July 2024. 2
- Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. [Progress measures for grokking via mechanistic interpretability](#). *ICLR*, January 2023a. 9, 21, 22, 23, 25, 29, 30
- Neel Nanda, Andrew Lee, and Martin Wattenberg. [Emergent linear representations in world models of self-supervised sequence models](#). *BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, September 2023b. 8, 11, 28