

- Neel Nanda, S. Rajamanoharan, J. Kramár, and R. Shah. [Fact finding: Attempting to reverse-engineer factual recall on the neuron level](#). *AI Alignment Forum*, 2023c. URL <https://www.alignmentforum.org/posts/iGuwZTHWb6DFY3sKB/fact-finding-attempting-to-reverse-engineer-factual-recall>, 2023c. 10
- Geraldin Nanfack, Alexander Fulleringer, Jonathan Marty, Michael Eickenberg, and Eugene Belilovsky. [Adversarial attacks on the interpretation of neuron activation maximization](#). *Proceedings of the AAAI Conference on Artificial Intelligence*, March 2024. 13
- Richard Ngo, Lawrence Chan, and Sören Mindermann. [The alignment problem from a deep learning perspective](#). *CoRR*, December 2022. 24
- Thanh Tam Nguyen, Thanh Trung Huynh, Phi Le Nguyen, Alan Wee-Chung Liew, Hongzhi Yin, and Quoc Viet Hung Nguyen. [A survey of machine unlearning](#). *CoRR*, October 2022. 24, 25
- Eshaan Nichani, Alex Damian, and Jason D. Lee. [How transformers learn causal structure with gradient descent](#). *CoRR*, February 2024. 11
- NicholasKees and janus. [Searching for search](#). *AI Alignment Forum*, November 2022. 24
- nostalgebraist. [interpreting gpt: the logit lens](#). *AI Alignment Forum*, August 2020. 13, 14, 19
- Chris Olah. [Distributed representations: Composition & superposition](#). *Transformer Circuits Thread*, 2023. 27
- Chris Olah and Shan Carter. [Research debt](#). *Distill*, March 2017. 2
- Chris Olah, Alexander Mordvintsev, and Ludwig Schubert. [Feature visualization](#). *Distill*, November 2017. 29
- Chris Olah, Arvind Satyanarayan, Ian Johnson, Shan Carter, Ludwig Schubert, Katherine Ye, and Alexander Mordvintsev. [The building blocks of interpretability](#). *Distill*, March 2018. 2
- Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter. [Zoom in: An introduction to circuits](#). *Distill*, March 2020. 2, 4, 5, 9, 10, 21, 32, 34
- Christopher Olah. [Mechanistic interpretability, variables, and the importance of interpretable bases](#). *Transformer Circuits Thread*, 2022. 2, 4
- B. A. Olshausen and D. J. Field. [Sparse coding with an overcomplete basis set: a strategy employed by v1?](#) *Vision Res*, December 1997. 15, 16
- Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. [In-context learning and induction heads](#). *Transformer Circuits Thread*, 2022. 10, 20, 22, 25, 30, 32
- Laura O’Mahony, Vincent Andrearczyk, Henning Muller, and Mara Graziani. [Disentangling neuron representations with concept vectors](#). *CVPR Workshops*, April 2023. 9
- Charles O’Neill and Thang Bui. [Sparse autoencoders enable scalable and reliable circuit identification in language models](#). *CoRR*, May 2024. 15, 19
- Francesco Ortu, Zhijing Jin, Diego Doimo, Mrinmaya Sachan, Alberto Cazzaniga, and Bernhard Schölkopf. [Competition of mechanisms: Tracing how language models handle facts and counterfactuals](#). *ACL 2024*, June 2024. 10
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. [Training language models to follow instructions with human feedback](#). *CoRR*, 2022. 30

- Koyena Pal, Jiuding Sun, Andrew Yuan, Byron C. Wallace, and David Bau. [Future lens: Anticipating subsequent tokens from a single hidden state](#). *CoNLL*, 2023. 15, 29
- Vedant Palit, Rohan Pandey, Aryaman Arora, and P. Liang. [Towards vision-language mechanistic interpretability: A causal tracing tool for blip](#). *ICCVW*, 2023. 30
- Xu Pan, Aaron Philip, Ziqian Xie, and Odelia Schwartz. [Dissecting query-key interaction in vision transformers](#). *ICML MI Workshop*, June 2024. 30
- Kiho Park, Yo Joong Choe, and Victor Veitch. [The linear representation hypothesis and the geometry of large language models](#). *NeurIPS Workshop on Causal Representation Learning*, November 2023a. 28
- Kiho Park, Yo Joong Choe, Yibo Jiang, and Victor Veitch. [The geometry of categorical and hierarchical concepts in large language models](#). *ICML MI Workshop (Oral)*, June 2024. 9
- Peter S. Park, Simon Goldstein, Aidan O’Gara, Michael Chen, and Dan Hendrycks. [Ai deception: A survey of examples, risks, and potential solutions](#). *CoRR*, August 2023b. 24
- Roma Patel and Ellie Pavlick. [Mapping language models to grounded conceptual spaces](#). *ICLR*, 2022. 12
- Judea Pearl. *Causality*. Cambridge University Press, 2009. 11, 16
- Judea Pearl and Dana Mackenzie. *The Book of Why: The New Science of Cause and Effect*. Penguin Books Limited, May 2018. 28
- Jonas Peters, Peter Bühlmann, and Nicolai Meinshausen. [Causal inference using invariant prediction: identification and confidence intervals](#). *Journal of the Royal Statistical Society Series B: Statistical Methodology*, November 2015. 28
- Jonas Peters, Dominik Janzing, and Bernhard Schölkopf. *Elements of Causal Inference: Foundations and Learning Algorithms*. The MIT Press, October 2017. 28
- Nicky Pochinkov and Nandi . [Machine unlearning evaluations as interpretability benchmarks](#). *AI Alignment Forum*, August 2023. 25
- Michael Poli, Stefano Massaroli, Eric Nguyen, Daniel Y. Fu, Tri Dao, Stephen Baccus, Yoshua Bengio, Stefano Ermon, and Christopher Ré. [Hyena hierarchy: Towards larger convolutional language models](#). *ICML*, April 2023. 25
- Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. [Grokking: Generalization beyond overfitting on small algorithmic datasets](#). *CoRR*, January 2022. 21
- Nikhil Prakash, Tamar Rott Shaham, Tal Haklay, Yonatan Belinkov, and David Bau. [Fine-tuning enhances existing mechanisms: A case study on entity tracking](#). *ICLR*, February 2024. 30
- Philip Quirke and Fazl Barez. [Understanding addition in transformers](#). *CoRR*, October 2023. 9, 29, 30
- Philip Quirke, Clement Neo, and Fazl Barez. [Increasing trust in language models through the reuse of verified circuits](#). *CoRR*, February 2024. 29
- Daking Rai, Yilun Zhou, Shi Feng, Abulhair Saparov, and Ziyu Yao. [A practical review of mechanistic interpretability for transformer-based language models](#). *CoRR*, July 2024. 2
- Senthoooran Rajamanoharan, Arthur Conmy, Lewis Smith, Tom Lieberum, Vikrant Varma, János Kramár, Rohin Shah, and Neel Nanda. [Improving dictionary learning with gated sparse autoencoders](#). *CoRR*, April 2024. 16
- Tilman Räuker, Anson Ho, Stephen Casper, and Dylan Hadfield-Menell. [Toward transparent ai: A survey on interpreting the inner structures of deep neural networks](#). *TMLR*, August 2023. 1, 13, 25, 26, 27, 28, 29

- Shauli Ravfogel, Yanai Elazar, Hila Gonen, Michael Twiton, and Yoav Goldberg. [Null it out: Guarding protected attributes by iterative nullspace projection](#). *ACL*, July 2020. 18
- Jie Ren, Mingjie Li, Qirui Chen, Huiqi Deng, and Quanshi Zhang. [Defining and quantifying the emergence of sparse concepts in dnns](#). *CoRR*, April 2023. 27
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. ["why should i trust you?": Explaining the predictions of any classifier](#). *NAACL*, August 2016. 2, 18
- RicG. [Agi-automated interpretability is suicide](#). *LessWrong*, May 2023. 25, 30
- Jonathan Richens and Tom Everitt. [Robust agents learn causal world models](#). *ICLR Oral*, February 2024. 11
- Ryan Riegel, Alexander Gray, Francois Luus, Naweed Khan, Ndivhuwo Makondo, Ismail Yunus Akhalwaya, Haifeng Qian, Ronald Fagin, Francisco Barahona, Udit Sharma, Shajith Ikbal, Hima Karanam, Sumit Neelam, Ankita Likhyan, and Santosh Srivastava. [Logical neural networks](#). *NeurIPS*, June 2020. 20, 27
- Mateo Rojas-Carulla, Bernhard Scholkopf, Richard Turner, and Jonas Peters. Invariant models for causal transfer learning. *JMLR*, 2018. 28
- Andrew Slavin Ross and Finale Doshi-Velez. [Improving the adversarial robustness and interpretability of deep neural networks by regularizing their input gradients](#). *AAAI*, November 2017. 28
- Cynthia Rudin. [Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead](#). *Nat Mach Intell*, May 2019. 29
- Cody Rushing and Neel Nanda. [Explorations of self-repair in language models](#). *ICML*, May 2024. 18
- Thane Ruthenis. [Internal interfaces are a high-priority interpretability target](#). *AI Alignment Forum*, December 2022. 24
- Thane Ruthenis. [World-model interpretability is all we need](#). *AI Alignment Forum*, January 2023. 24
- Hassan Sajjad, Nadir Durrani, and Fahim Dalvi. [Neuron-level interpretation of deep nlp models: A survey](#). *TACL*, November 2022. 5
- Mansi Sakarvadia, Aswathy Ajith, Arham Khan, Daniel Grzenda, Nathaniel Hudson, André Bauer, Kyle Chard, and Ian Foster. [Memory injections: Correcting multi-hop reasoning failures during inference in transformer-based language models](#). *CoRR*, September 2023a. 9
- Mansi Sakarvadia, Arham Khan, Aswathy Ajith, Daniel Grzenda, Nathaniel Hudson, André Bauer, Kyle Chard, and Ian Foster. [Attention lens: A tool for mechanistically interpreting the attention head information retrieval mechanism](#). *CoRR*, October 2023b. 15
- Emmanuelle Salin, Badreddine Farah, Stéphane Ayache, and Benoit Favre. [Are vision-language transformers learning multimodal representations? a probing perspective](#). *AAAI*, June 2022. 30
- Tommaso Salvatori, Ankur Mali, Christopher L. Buckley, Thomas Lukasiewicz, Rajesh P. N. Rao, Karl Friston, and Alexander Ororbia. [Brain-inspired computational intelligence via predictive coding](#). *CoRR*, 2023. 11
- Naomi Saphra. [Interpretability creationism](#). *The Gradient*, 2023. 21
- Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. [Are emergent abilities of large language models a mirage?](#) *CoRR*, May 2023. 20
- Adam Scherlis, Kshitij Sachan, Adam S. Jermyn, Joe Benton, and Buck Shlegeris. [Polysemanticity and capacity in neural networks](#). *CoRR*, July 2023. 6, 26, 27