

- Bernhard Schölkopf, Francesco Locatello, Stefan Bauer, Nan Rosemary Ke, Nal Kalchbrenner, Anirudh Goyal, and Yoshua Bengio. [Towards causal representation learning](#). *Special Issue of Proceedings of the IEEE - Advances in Machine Learning and Deep Neural Networks*, February 2021. 27
- Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Metzler. [Confident adaptive language modeling](#). *NeurIPS Oral*, October 2022. 25
- Ramprasaath R. Selvaraju, Abhishek Das, Ramakrishna Vedantam, Michael Cogswell, Devi Parikh, and Dhruv Batra. [Grad-cam: Why did you say that? visual explanations from deep networks via gradient-based localization](#). *ICCV*, 2016. 2
- Murray Shanahan, Kyle McDonell, and Laria Reynolds. [Role play with large language models](#). *Nature*, November 2023. 11, 12, 24
- Lloyd S. Shapley. [A value for \$n\$ -person games](#). *Cambridge University Press*, October 1988. 2
- Lee Sharkey. [Circumventing interpretability: How to defeat mind-readers](#). *CoRR*, December 2022. 24, 26
- Lee Sharkey. [A technical note on bilinear layers for interpretability](#). *CoRR*, May 2023. 20
- Lee Sharkey, Sid Black, and beren. [Current themes in mechanistic interpretability research](#). *AI Alignment Forum*, November 2022a. 2, 8
- Lee Sharkey, Dan Braun, and Beren Millidge. [Taking features out of superposition with sparse autoencoders](#). *AI Alignment Forum*, 2022b. 8, 15, 16
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. [Towards understanding sycophancy in language models](#). *CoRR*, October 2023. 28, 35
- Justin Shovelain and Elliot McKernon. [The risk-reward tradeoff of interpretability research](#). *LessWrong*, July 2023. 25
- Avanti Shrikumar, Peyton Greenside, and Anshul Kundaje. [Learning important features through propagating activation differences](#). *ICML*, 2017. 2
- James B. Simon, Maksis Knutins, Liu Ziyin, Daniel Geisz, Abraham J. Fetterman, and Joshua Albrecht. [On the stepwise nature of self-supervised learning](#). *ICML*, May 2023. 20
- slavachalnev. [Sparse mlp distillation](#). *LessWrong*, 2024. 16
- Daniel Smilkov, Nikhil Thorat, Been Kim, Fernanda Viégas, and Martin Wattenberg. [Smoothgrad: removing noise by adding noise](#). *CoRR*, June 2017. 2
- Nate Soares. [If interpretability research goes well, it may get dangerous](#). *LessWrong*, April 2023. 25
- Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. [Dropout: A simple way to prevent neural networks from overfitting](#). *JMLR*, 2014. 5
- Dashiell Stander, Qinan Yu, Honglu Fan, and Stella Biderman. [Grokking group multiplication with cosets](#). *CoRR*, 2023. 21
- Jacob Steinhardt. [Emergent deception and emergent optimization](#). *Bounded Regret*, February 2023. 20, 23, 28
- Alessandro Stolfo, Yonatan Belinkov, and Mrinmaya Sachan. [A mechanistic interpretation of arithmetic reasoning in language models using causal mediation analysis](#). *EMNLP*, October 2023. 17, 22

- Ilya Sucholutsky, Lukas Muttenthaler, Adrian Weller, Andi Peng, Andreea Bobu, Been Kim, Bradley C. Love, Erin Grant, Iris Groen, Jascha Achterberg, Joshua B. Tenenbaum, Katherine M. Collins, Katherine L. Hermann, Kerem Oktar, Klaus Greff, Martin N. Hebart, Nori Jacoby, Qiuyi Zhang, Raja Marjeh, Robert Geirhos, Sherol Chen, Simon Kornblith, Sunayana Rane, Talia Konkle, Thomas P. O’Connell, Thomas Unterthiner, Andrew K. Lampinen, Klaus-Robert Müller, Mariya Toneva, and Thomas L. Griffiths. [Getting aligned on representational alignment](#). *CoRR*, November 2023. 10
- Chen Sun, Nolan Andrew Miller, Andrey Zhmoginov, Max Vladymyrov, and Mark Sandler. [Learning and unlearning of fabricated knowledge in language models](#). *ICML MI Workshop*, June 2024. 25
- Mukund Sundararajan, Ankur Taly, and Qiqi Yan. [Axiomatic attribution for deep networks](#). *ICML*, June 2017. 2
- Aaquib Syed, Can Rager, and Arthur Conmy. [Attribution patching outperforms automated circuit discovery](#). *CoRR*, October 2023. 18, 19, 22
- technicalities and Stag. [Shallow review of live agendas in alignment & safety](#). *LessWrong*, 2023. 24
- Max Tegmark and Steve Omohundro. [Provably safe systems: the only path to controllable agi](#). *CoRR*, September 2023. 24, 29, 35
- Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, and Brian Chen. [Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet](#). *Transformer Circuits Thread*, 2024. 15
- Ian Tenney, Dipanjan Das, and Ellie Pavlick. [Bert rediscovered the classical nlp pipeline](#). *ACL*, August 2019. 14
- Vimal Thilak, Etai Littwin, Shuangfei Zhai, Omid Saremi, Roni Paiss, and Joshua Susskind. [The slingshot mechanism: An empirical study of adaptive optimizers and the grokking phenomenon](#). *CoRR*, 2022. 21
- Hannes Thurnherr and Jérémy Scheurer. [Tracrbench: Generating interpretability testbeds with large language models](#). *ICML MI Workshop*, June 2024. 29
- Curt Tigges, Oskar John Hollinsworth, Atticus Geiger, and Neel Nanda. [Language models linearly represent sentiment](#). *ICML MI Workshop*, June 2024. 9
- Eric Todd, Millicent L. Li, Arnab Sen Sharma, Aaron Mueller, Byron C. Wallace, and David Bau. [Function vectors in large language models](#). *CoRR*, 2023. 9
- Dweep Trivedi, Jesse Zhang, Shao-Hua Sun, and Joseph J. Lim. [Learning to synthesize programs as interpretable and generalizable policies](#). *NeurIPS*, 2021. 27
- Alexander Matt Turner, Lisa Thiergart, David Udell, Gavin Leech, Ulisse Mini, and Monte MacDiarmid. [Activation addition: Steering language models without optimization](#). *CoRR*, September 2023. 9
- Dmitry Vaintrob, jake_mendel, and Kaarel. [Toward a mathematical framework for computation in superposition](#). *AI Alignment Forum*, 2024. 29
- Alexandre Variengien and Eric Winsor. [Look before you leap: A universal emergent decomposition of retrieval tasks in language models](#). *ICML MI Workshop*, December 2023. 11, 22, 30
- Vikrant Varma, Rohin Shah, Zachary Kenton, János Kramár, and Ramana Kumar. [Explaining grokking through circuit efficiency](#). *CoRR*, September 2023. 21, 22
- Abhinav Verma, Hoang M. Le, Yisong Yue, and Swarat Chaudhuri. [Imitation-projected programmatic reinforcement learning](#). *NeurIPS*, 2019a. 27
- Abhinav Verma, Vijayaraghavan Murali, Rishabh Singh, Pushmeet Kohli, and Swarat Chaudhuri. [Programmatically interpretable reinforcement learning](#). *CoRR*, April 2019b. 27

- Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. [Investigating gender bias in language models using causal mediation analysis](#). *NeurIPS*, 2020. 17, 18
- M. Vilas, Timothy Schaumlöffel, and Gemma Roig. [Analyzing vision transformers for image classification in class embedding space](#). *CoRR*, 2023. 30
- Elena Voita and Ivan Titov. [Information-theoretic probing with minimum description length](#). *EMNLP*, March 2020. 14
- Elena Voita, Javier Ferrando, and Christoforos Nalmpantis. [Neurons in large language models: Dead, n-gram, positional](#). *CoRR*, September 2023. 5
- Julius von Kügelgen, M. Loog, A. Mey, and B. Scholkopf. [Semi-supervised learning, causality, and the conditional cluster assumption](#). *Conference on Uncertainty in Artificial Intelligence*, May 2019. 28
- Johannes von Oswald, Eyvind Niklasson, Maximilian Schlegel, Seijin Kobayashi, Nicolas Zucchet, Nino Scherrer, Nolan Miller, Mark Sandler, Blaise Agüera y Arcas, Max Vladymyrov, Razvan Pascanu, and João Sacramento. [Uncovering mesa-optimization algorithms in transformers](#). *CoRR*, September 2023. 24, 26
- Chelsea Voss, Gabriel Goh, Nick Cammarata, Michael Petrov, Ludwig Schubert, and Chris Olah. [Branch specialization](#). *Distill*, April 2021. 10, 21
- Boshi Wang, Xiang Yue, Yu Su, and Huan Sun. [Grokked transformers are implicit reasoners: A mechanistic journey to the edge of generalization](#). *ICML MI Workshop*, June 2024. 21
- Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. [Interpretability in the wild: a circuit for indirect object identification in gpt-2 small](#). *ICLR*, 2023. 10, 13, 17, 18, 19, 22, 30, 32
- Alex Warstadt, Alicia Parrish, Haokun Liu, Anhad Mohananey, Wei Peng, Sheng-Fu Wang, and Samuel R. Bowman. [Blimp: The benchmark of linguistic minimal pairs for english](#). *Transactions of the Association for Computational Linguistics*, 2020. 2
- Sumio Watanabe. *Algebraic Geometry and Statistical Learning Theory*. Cambridge Monographs on Applied and Computational Mathematics. Cambridge University Press, 2009. 21
- Sumio Watanabe. *Mathematical Theory of Bayesian Statistics*. Chapman and Hall, 1 edition, April 2018. 21
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. [Emergent abilities of large language models](#). *TMLR*, October 2022. 20, 23
- John Wentworth. [How to go from interpretability to alignment: Just retarget the search](#). *AI Alignment Forum*, August 2022. 24
- James C. R. Whittington, Will Dorrell, Surya Ganguli, and Timothy E. J. Behrens. [Disentangling with biological constraints: A theory of functional cell types](#). *CoRR*, September 2022. 15
- Baoyuan Wu, Hongrui Chen, Mingda Zhang, Zihao Zhu, Shaokui Wei, Danni Yuan, and Chao Shen. [Backdoorbench: A comprehensive benchmark of backdoor learning](#). *NeurIPS Datasets and Benchmarks*, October 2022a. 29
- Zhengxuan Wu, Atticus Geiger, Joshua Rozner, Elisa Kreiss, Hanson Lu, Thomas Icard, Christopher Potts, and Noah Goodman. [Causal distillation for language models](#). *NAACL-HLT*, July 2022b. 19